

RESEARCH MEMORANDUM

ASIC-DIFFUSION: Hardware-Accelerated Latent Space Generation for Diffusion Models Using Repurposed Mining ASICs

Document ID:	ASIC-DIFF-2025-001
Principal Investigator:	Francisco Angulo de Lafuente
Date:	December 2024
Research Line:	Artistic Image Generation via Hardware Diffusion
Hardware Platform:	Lucky Miner LV06 (SHA-256 ASIC)

1. Executive Summary

This memorandum establishes the research framework for ASIC-DIFFUSION, an experimental program investigating the use of obsolete cryptocurrency mining hardware as physical latent space generators for diffusion-based image synthesis models. The core hypothesis posits that the deterministic chaos properties of SHA-256 ASIC hardware can provide computationally efficient, cryptographically-secured initialization states for denoising diffusion probabilistic models (DDPMs), potentially reducing the energy footprint and computational overhead of the initial noise generation phase by orders of magnitude.

The research addresses a fundamental inefficiency in current diffusion architectures: the generation of high-dimensional Gaussian noise tensors through pseudo-random number generators (PRNGs) implemented in software on general-purpose GPUs. By replacing software-generated stochasticity with hardware-derived deterministic chaos, we aim to demonstrate both economic and energetic advantages while potentially discovering novel aesthetic properties emergent from the physical characteristics of silicon-based hash computation.

2. Problem Statement

2.1 The Inefficiency of Software-Generated Latent Spaces

Contemporary diffusion models (Stable Diffusion, DALL-E 3, Midjourney) require initialization from a high-entropy latent tensor, typically sampled from a standard Gaussian distribution $N(0,1)$. This sampling operation, while mathematically trivial, presents several engineering challenges at scale:

- **Energy Overhead:** GPU-based PRNGs consume significant power for entropy generation, estimated at 0.5-2W per sampling operation on modern architectures.
- **Latency Contribution:** Initial tensor generation contributes 50-150ms to total inference time, representing 1-5% of the generation pipeline.
- **Reproducibility Complexity:** Seed-based reproduction requires software state management and is vulnerable to floating-point precision variations across hardware platforms.
- **Scalability Bottleneck:** Each concurrent user requires dedicated GPU resources for noise generation, limiting horizontal scaling efficiency.

2.2 The ASIC Opportunity

The global cryptocurrency mining industry has produced millions of Application-Specific Integrated Circuits (ASICs) optimized for SHA-256 hash computation. As mining difficulty increases, older hardware becomes economically unviable for its original purpose, creating a

secondary market of high-performance, low-cost silicon. The Lucky Miner LV06, our target platform, represents this category: a 500 GH/s SHA-256 accelerator consuming only 3.5W, available for under \$30 USD.

The SHA-256 algorithm exhibits the *avalanche effect*: a single-bit change in input produces a statistically independent output with approximately 50% bit-flip probability. This property, combined with the deterministic nature of hash functions, creates a unique computational primitive: **reproducible chaos**.

3. Research Objectives

3.1 Primary Objectives

1. **Validate Hardware-Software Integration:** Demonstrate reliable communication between the LV06 ASIC and standard computing infrastructure via the Stratum mining protocol, achieving sustained hash generation rates exceeding 100 GH/s.
2. **Characterize Output Distribution:** Statistically analyze the bit distribution, autocorrelation, and spectral properties of LV06-generated hashes to determine suitability as latent space initializers.
3. **Implement Latent Space Mapping:** Develop transformation functions to convert 256-bit SHA-256 outputs into latent tensors compatible with Stable Diffusion v1.5 ($512 \times 512 \times 4$) and SDXL ($1024 \times 1024 \times 4$) architectures.
4. **Quantify Efficiency Gains:** Measure energy consumption (Joules per latent tensor) and throughput (tensors per second) compared to GPU-based PRNG implementations.

3.2 Secondary Objectives

- Investigate emergent aesthetic properties of ASIC-initialized diffusion outputs.
- Develop open-source tooling for ASIC-diffusion integration.
- Explore multi-ASIC parallelization for high-throughput inference servers.
- Document reproducibility protocols for cryptographically-verifiable image generation.

4. Technical Architecture

4.1 Hardware Layer

The Lucky Miner LV06 serves as the physical entropy source. Key specifications include 500 GH/s SHA-256 hashrate, 3.5W power consumption, USB 2.0 connectivity, and BM1387 ASIC chipset compatibility. The device communicates via a modified Stratum protocol, accepting work units containing 80-byte block headers and returning 4-byte nonce solutions.

4.2 Signal Layer

A Python-based middleware translates between the ASIC's mining protocol and the diffusion model's tensor requirements. The signal layer performs nonce injection by encoding seed values into the block header's nonce field, hash collection by capturing computed SHA-256 outputs at wire speed, and tensor assembly by aggregating sequential hashes into latent-compatible matrices.

4.3 Readout Layer

The readout layer transforms raw hash outputs into diffusion-compatible latent representations. Two approaches will be evaluated: direct mapping, which interprets hash bits as floating-point values in $[-1, 1]$ range via linear scaling, and learned projection, which trains a small neural network to map hash space to optimal latent distributions.

5. Experimental Phases

1. **Phase 1 - Hardware Validation (Weeks 1-2):** Establish stable communication with LV06, verify hashrate consistency, implement basic data collection infrastructure.
2. **Phase 2 - Statistical Characterization (Weeks 3-4):** Generate 10 million+ hashes, perform comprehensive statistical analysis, validate randomness properties against NIST SP 800-22 test suite.
3. **Phase 3 - Integration Development (Weeks 5-8):** Implement Stable Diffusion integration, develop transformation pipelines, create benchmark infrastructure.
4. **Phase 4 - Comparative Evaluation (Weeks 9-12):** Energy efficiency measurements, throughput benchmarking, image quality assessment via FID and CLIP scores.
5. **Phase 5 - Aesthetic Exploration (Weeks 13-16):** Investigate emergent visual properties, document unique artifacts, develop artistic applications.

6. Expected Outcomes

Upon successful completion, this research program will deliver quantified efficiency metrics demonstrating energy savings in latent space generation, open-source software toolkit for ASIC-diffusion integration, comprehensive statistical characterization of hardware-generated entropy, comparative analysis with GPU-based approaches across multiple dimensions, and documentation of novel aesthetic properties emergent from hardware-based initialization.

7. Resource Requirements

Hardware	Software
Lucky Miner LV06 × 1	Python 3.10+
USB power meter (precision)	PyTorch / Diffusers library
Host computer (any modern CPU)	Custom Stratum client
Optional: GPU for diffusion inference	Statistical analysis suite (NumPy, SciPy)

8. Conclusion

The ASIC-DIFFUSION research program represents a novel approach to democratizing AI image generation by repurposing abundant, low-cost mining hardware. Beyond immediate efficiency gains, this research explores fundamental questions about the nature of computational creativity: whether the physical properties of silicon can introduce aesthetically meaningful variations into generative processes. Success would establish a new paradigm for sustainable, accessible AI infrastructure built on electronic waste streams rather than premium semiconductor manufacturing.

Author Information

Name: Francisco Angulo de Lafuente
Affiliation: Independent Researcher

GitHub: <https://github.com/Agnuxo1>
ResearchGate: <https://www.researchgate.net/profile/Francisco-Angulo-Lafuente-3>
Kaggle: <https://www.kaggle.com/franciscoangulo>
HuggingFace: <https://huggingface.co/Agnuxo>
Wikipedia: https://es.wikipedia.org/wiki/Francisco_Angulo_de_Lafuente